

Genetic Algorithms

1. Genetics – and genetic algorithms
2. The travelling salesman
3. Resource optimisation problems
4. Clustering
5. Rule generation

Genetics and genetic algorithms

- Nature has found effective ways to solve many problems – we can mimic its approaches in “biologically-inspired” computer systems
 - Neural networks
 - Genetic algorithms (GAs)
 - Artificial life
 - Swarm intelligence


Optimisation problems

- Optimisation problems are about finding the best ...
- e.g.,
 - The best allocation of classes to rooms
 - The best clustering
 - The best prediction of ...
- Many data mining problems can be rephrased as optimisation problems

Optimisation problems


- Are phrased as:
 - A set of parameters (input variables)
 - A fitness function defined on these parameters (which measures “goodness”)
 - A set of constraints (on the parameters)
- The problem is then to find a (or the) set of parameter values that satisfy the constraints & maximise the fitness function

Optimisation and hill climbing




A hill climbing attempt to find the global maxima (i.e., to find a) can easily get stuck at the local maxima – because at each step it only makes a local analysis

A parallel attempt using hill climbing



Because of the parallelism, at least some of the partial solutions will (hopefully) reach the global maxima – even though some others get stuck at the local maxima

A parallel attempt using natural selection



Evolution/natural selection creates new individuals from old by breeding, the hope being that “good” individuals will have good children; less good individuals have a lower chance of survival (and hence of breeding)

Genetics

- DNA consists of a sequence of nucleotides, each being based on one of 4 sub-bases
 - adenine (A); cytosine (C); guanine (G); and thymine (T)
- Triples of the above represent amino acids
 - e.g., AAG represents Lysine
- A gene is a portion of the DNA string that
 - encodes a particular function
 - is the unit of heredity


Genetics

- *Mutation* is where one of the genes is modified slightly in some random fashion
- *Crossover* (breeding) is the construction of a new DNA string – with some of the genes coming from the first parent and the rest coming from the other parent

Genetic algorithms

- Code up the problem as a DNA string
 - Each potential/partial solution is coded up as a particular DNA string (i.e., as an individual in our artificial world)
 - This coding is the creative bit
- Define fitness function
 - how good an individual DNA string is
- Define crossover and mutation operators
- Define initial population ... and run: in each generation allow some of the individuals to breed/mutate and remove some of the less good individuals
- When we reach stability – i.e., we don't seem to be improving our goodness of fit, . . .

2. The travelling salesman



■ = office ● = site to be visited

The travelling salesman problem

- For each pair of sites, there is a distance between them. The problem is simple:
 - what's the shortest route?
- Brute-force search: exponential complexity
- Important because:
 - A very wide class of computational problems are known to be equivalent to this problem
 - Unknown if there is any polynomial-time algorithm ... general belief is “no”


Encoding the travelling salesman

- Number the sites
- Each potential solution is simply a sequence of site numbers
 - e.g., with 8 sites: (3,8,7,1,5,2,6,4)
- Fitness function is travel distance in the given order (low distance = high fitness)

Encoding the travelling salesman

- Mutation
 - Each value must appear once in the DNA string, therefore modifying one of the values does not make sense
 - Therefore mutation consists of interchanging two values

$(\underline{3}, 8, 7, \underline{1}, 5, 2, 6, 4) \rightarrow (1, 8, 7, 3, 5, 2, 6, 4)$



Encoding the travelling salesman

- Crossover

- Each value must appear once in the DNA string: some crossovers will not “make sense”
- Given 2 individuals with a “common initial segment”, take the initial segment of one with the latter part of the other

(3,8,7,1,5,4,6,2) & (8,7,3,2,5,1,6,4)



(3,8,7,2,5,1,6,4)


3. Resource allocation

- Consider the timetabling problem
- There are a set of classes and rooms
- Constraints and preferences are obvious
 - Rooms should be bigger than classes
 - No double booking
 - All required classes are timetabled
 - Students & teachers should have at most X hours/day
 - Other teacher preferences

DNA construction

- All DNA strings are an array containing an entry for each “room slot”
 - Room; Date; Time
- The *value in* the room slot is the name of the class allocated
- Fitness can be defined in terms of the violation of preferences (soft constraints)

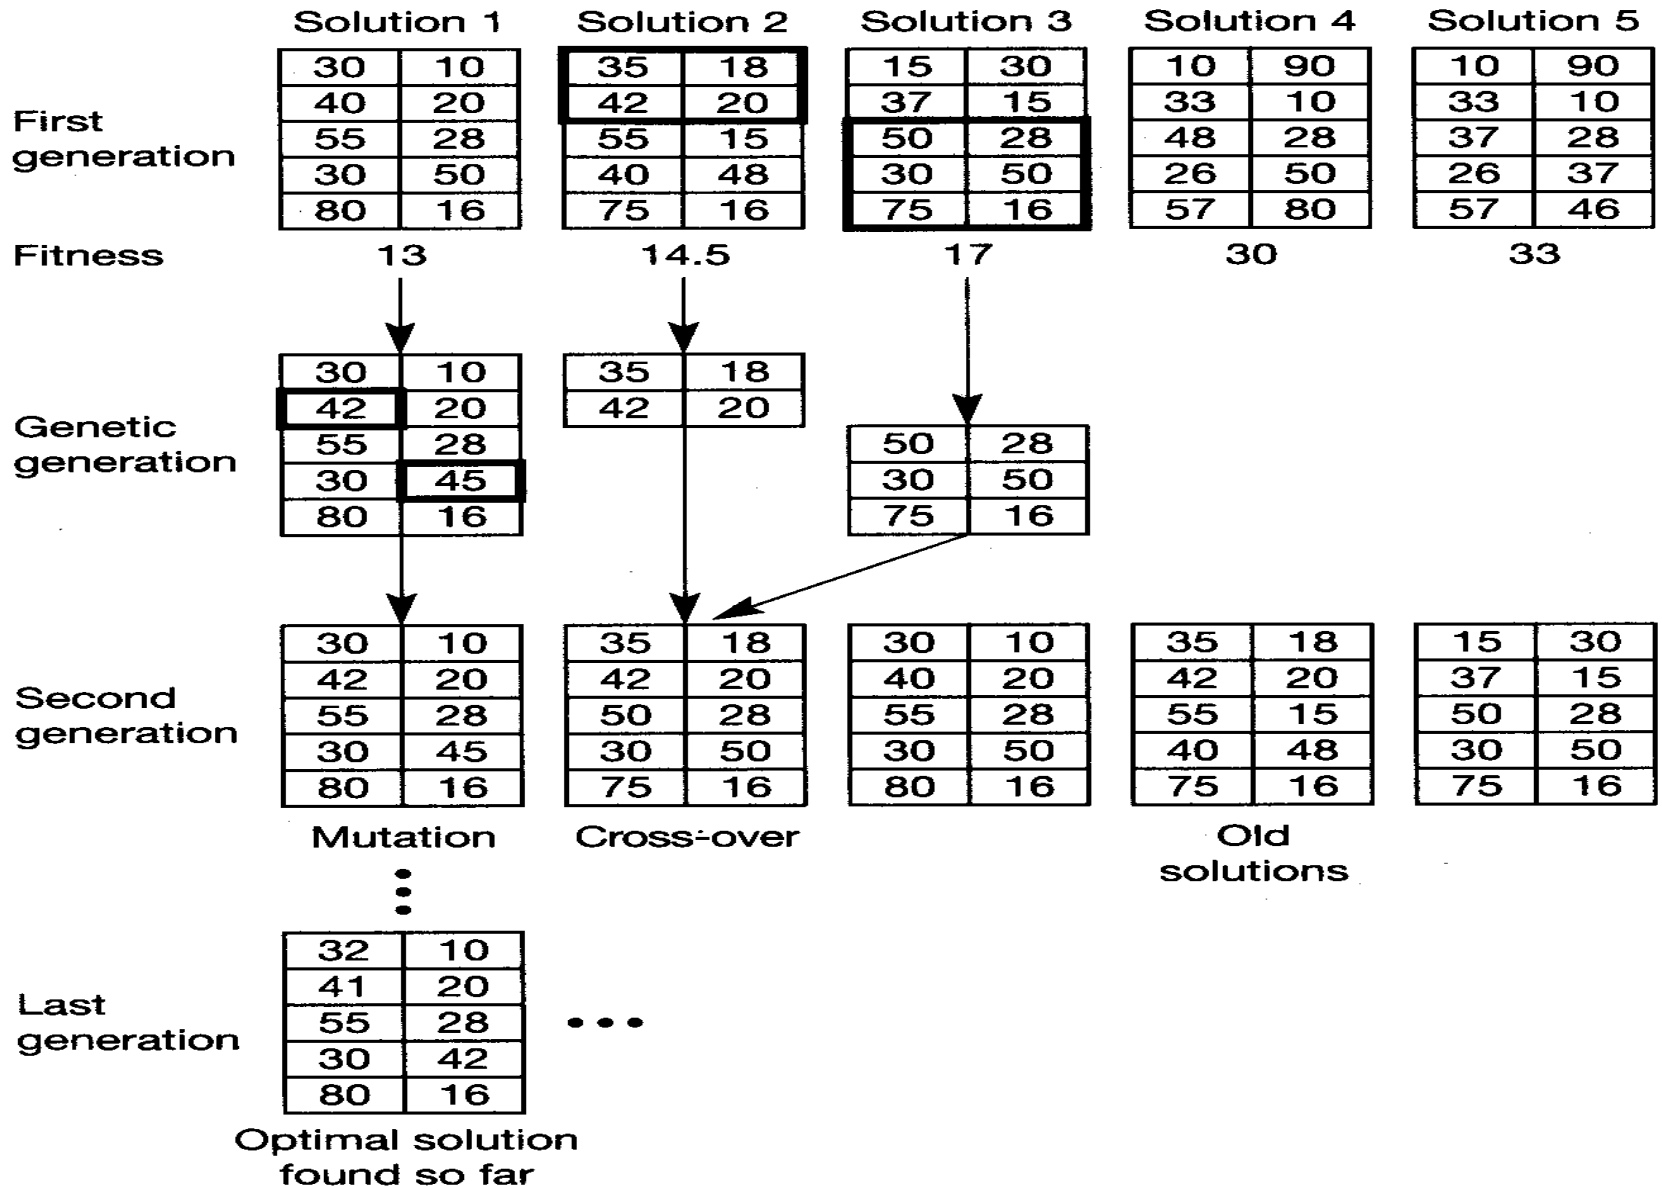
4. Clustering



Genetic Algorithms

- Fitness function as before
 - internal cohesion, external separation
- Cross over: take 3 points in one sequence and 2 in another - combine
- Mutation: slightly perturb one or more figures

Initial set of random solutions



5. Rule generation

- For classification, a rule has the form
$$\text{Cond}_1 \ \& \ \text{Cond}_2 \ \& \ \dots \ \& \ \text{Cond}_N \rightarrow \text{Class}=\text{C}$$
- Conditions are typically of the form
$$A=v; \ A \neq v; \ A > v; \ A < v$$
- Whatever the type of conditions allowed, the DNA string is obvious:
$$(\text{C}, \text{Cond}_1, \text{Cond}_2, \dots, \text{Cond}_N, \text{null}, \text{null}, \dots)$$

Rule generation

- The fitness of a DNA string/rule
 $\text{Cond}_1 \ \& \ \text{Cond}_2 \ \& \ ... \ \& \ \text{Cond}_N \rightarrow \text{Class}=\text{C}$
is clearly its predictive accuracy
- Mutation is straightforward
- Crossover is as previously defined – take some elements of one parent and some elements of the other